

COSC 407

Intro to Parallel Computing

Topic 7 - Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Outline

Previously:

- Synchronization (barriers, nowait)
- Mutual Exclusion (critical, atomic, locks)

Today

- **Variable scope (shared, private, firstprivate)**
- **Reduction**

Variable Scope


- In **serial programming**, the scope of a variable consists of those parts of a program in which the variable can be used
- In **OpenMP**, the scope of a variable refers to the **set of threads that can access the variable in a parallel block**
- A **shared variable** exists in only **one memory** location and all threads in the team access this location
 - All variables **declared BEFORE a parallel block** are shared by default
 - **shared(x)**
 - x will refer to the same memory block for all threads

Variable Scope


- A **private variable** can only be accessed by a single thread (**each thread has its own copy**).
 - variables **declared WITHIN a parallel block** are private by default
 - **private(y)**
 - y will refer to a different memory block for each thread. Each copy of y is *uninitialized*.
 - **firstprivate(z)** same as private, but each copy of z is *initialized* with the value that the original z has when the construct is encountered

Variable Scope, *cont'd*

```
int x = 5;  
#pragma omp parallel shared(x)
```



```
int y = 5;  
#pragma omp parallel private(y)  
//each thread creates a new copy  
of y, and these y's are uninitialized
```




Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Variable Scope, *cont'd*

```
int z = 5;  
#pragma omp parallel firstprivate(z)  
//each thread creates a new copy  
of z, and these z's are initialized  
with value of original z
```



- In some cases, the variables scope is predetermined and cannot be changed:
 - Variables declared inside the parallel region are private
 - A loop variable in a parallel loop is private
 - const variables are shared

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Variable Scope: Example

```
#include <omp.h>
#include <stdio.h>
int main() {
    int j=0, i;
    #pragma omp parallel private(i)           //i, j is shared by default
                                              //i is private in this block
    {
        printf("Started T%d\n", omp_get_thread_num());
        for (i = 0; i < 10000; i++)
            j++;
        printf("Finished T%d\n", omp_get_thread_num());
    }
    printf("%d\n", j);
    return 0;
}
```

Possible outputs
with 3 threads

Started T0
Finished T0
Started T1
Finished T1
Started T2
Finished T2
30000

Started T0
Started T1
Finished T1
Started T2
Finished T2
Finished T0
24624

Started T2
Started T0
Finished T0
Finished T2
Started T1
Finished T1
19616

race condition
(WHY?)

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

default clause

- Sets the default scope.

default(shared | private | none)

default(none) forces the programmer to specify the scope of each variable in a block – i.e. the compiler will require that we specify the scope of each variable we use in the block and that has been declared outside the block.

We don't have to mention z in #pragma as it is not used in the parallel region

```
int x = 0, y = 0, z = 0;
#pragma omp parallel num_threads(4) default(none) private(x) shared(y)
{
    x = omp_get_thread_num();
    #pragma omp atomic
    y = y + x;
}
printf("x:%d y:%d z:%d", x, y, z);
```

Try removing shared(y) or private(x) and notice what happens.

Output:
x:0 y:6 z:0

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Back to Sums.....

- Assume you want to find the sum of a function values in a range

```
double global_sum = 0;           // 1) create a shared global_sum
# pragma omp parallel num_threads(4)
{
    int my_id = ..;
    int my_sum = f(my_id); // 2) create a private my_sum and compute it
    #pragma omp critical
    global_sum += my_sum; // 3) update global sum in critical section
}
```

- Is there a better solution?

Reduction

- The reduction clause of parallel does the following:
 - Provides a private copy of a the variable for each thread
 - The variable is scoped 'reduction' (private first then shared on exit)
 - i.e., **no need for critical clause**
 - Combines the private results on exit

Reduction

Syntax:

`reduction (<op> : <variable list>)`

- A reduction operator is a binary operation: +, *, -, &, |, ^, &&, ||
 - **No division** (/)
 - op could also be **max** or **min**, (more later on this)
 - Assuming we have N threads and a variable x, on exit x's value is:
`x = init_value <op> x0 <op> x1 <op> ... <op> xN-1`
Where init_value is x's value, and x₀, x₁, x₂, etc are its private copies.
- **Initial values** for the temporary private variables:
`1` for (*, &&) `0` for (+, -, |, ^, ||) `~0` for (&)

Reduction: Example 1

```
int main() {
    int x = 10;           //shared x

    #pragma omp parallel reduction(+:x)
    {
        x = omp_get_thread_num(); //private x
        printf("Private x = %d\n", x);
    } //on exit: shared x += all private x's

    printf("Shared x = %d\n", x);
    return 0;
}
```

Possible output:

```
Private x = 3
Private x = 1
Private x = 0
Private x = 2
Shared x = 16
```

Reduction: Example 2

```
int main() {
    int x = 10, y = 10;           //shared x, y
    #pragma omp parallel reduction(+:x, y)
    {
        x = omp_get_thread_num(); //private x,y
        y = 5;
        printf("Private: x=%d, y=%d\n", x, y);
    } //on exit: shared x += all private x's
    // shared y += all private y's
    printf("Shared: x=%d,y=%d\n", x, y);
    return 0;
}
```

Possible output:

```
Private: x=1, y=5
Private: x=2, y=5
Private: x=0, y=5
Private: x=3, y=5
Shared: x=16,y=30
```

Caution

- Reduction using (-) is same as (+), so if you want to subtract
 - i.e. both `reduction(-:x)` and `reduction(+:x)` have the same effect; i.e. $\text{shared_x} = \text{init_value} + x_0 + x_1 + \dots$
- To do reduction on subtraction, use this code:

```
x = init_value();
#pragma omp parallel reduction(-:x)
x -= f(...);
```

- Why this works?
 - Initial x is 0, so value of private x is $-f()$
 - This means, final x is:
$$x = \text{init_value}() - f_0(\dots) - f_1(\dots) - \text{etc}$$
where $f_n()$ is the function value computed by thread id n

Back to the Sums...

```
double global_sum = 0;           // 1) create shared global_sum
# pragma omp parallel num_threads(4)
{
    int my_id = ..;
    int my_sum = f(my_id);       // 2) create a private my_sum
    #pragma omp critical
    global_sum += my_sum;        // 3) update global sum in critical section
}
```



```
double global_sum = 0;           //global_sum is shared
# pragma omp parallel num_threads(4) reduction(+:global_sum)
{
    int my_id = ..;
    global_sum = f(my_id); // each thread gets a private copy of global_sum
} //reduction is applied on exit
```


Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Area Under a Curve

```
int main() {
    double global_result = 0.0; // result stored in global_result
    double a = 1, b = 2;       // endpoints
    int thread_count = 4;      // should be = number of cores
    int n = 8;                 // Total number of trapezoids
                                // = multiple of thread_count

    # pragma omp parallel num_threads(thread_count)
    Trap(a, b, n, &global_result);
    printf("Approximate area: %f\n", global_result); return 0;
}
```



Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing


Area Under a Curve

```
void Trap(double a, double b, int n, double* global_result_p) {
    double h, x, my_result, my_a, my_b; int i, my_n;
    int my_id = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    h = (b - a) / n;
    my_n = n / thread_count; // # of contiguous trapezoids per thread
    my_a = a + (my_id * my_n) * h;
    my_b = my_a + my_n * h;

    my_result = (f(my_a) + f(my_b)) / 2.0;
    for (i = 1; i <= my_n - 1; i++) {
        x = my_a + i * h;
        my_result += f(x);
    }
    my_result = my_result * h;

    # pragma omp critical
    *global_result_p += my_result;
}
```



Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Area Calculation – v.2

small update without using reduction clause

- For the Trap function, instead of:

```
void Trap(double a, double b, int n, double* global_result_p)
```

we would prefer the more attractive version pf

```
double Local_trap(double a, double b, int n)
```

which...

- is run by each thread to return a part of the calculations
- has no critical section but....

```
double global_result = 0.0; //global_result is shared
# pragma omp parallel num_threads(thread_count)
{
    # pragma omp critical ← Why?
    global_result += Local_trap(a, b, n);
}
```

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Area Calculation – v.2

(small update without using reduction clause)

```
main() {
    double global_result = 0.0, a = 1, b = 2;
    int n = 12;
    # pragma omp parallel num_threads(4)
    {
        # pragma omp critical
        global_result += Local_trap(a, b, n); //global_result is shared
    }
    printf("Approximate area: %f\n", global_result); return 0;
}
```

Warning: sequential execution!
Q1) Explain! Q2) How to fix?

```
double Local_trap(double a, double b, int n) {
    double h, x, my_result, local_a, local_b;
    int i, local_n, my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    h = (b - a) / n;
    local_n = n / thread_count;
    local_a = a + my_rank * local_n * h;
    local_b = local_a + local_n * h;
    my_result = (f(local_a) + f(local_b)) / 2.0;
    for (i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    return h*my_result; //instead of adding it to global_result
}
```

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Area Calculation – v.3

(still not using reduction() clause)

- We can avoid the problem (of sequentially running the program) by declaring a private variable inside the parallel block and moving the critical section after the function call.

```
double global_result = 0.0; //global_result is shared
# pragma omp parallel num_threads(thread_count)
{
    double my_result = Local_trap(a,b,n); //my_result is private
    # pragma omp critical
    global_result += my_result;
}
```

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Area Calculation – v.3

(still not using `reduction()` clause)

```
main() {
    double global_result = 0.0, a=1, b=2;    //global_result is shared
    int n = 12, thread_count = 4;
    # pragma omp parallel num_threads(thread_count)
    {
        double my_result = Local_trap(a, b, n); //my_result is private
        # pragma omp critical
            global_result += my_result;
    }
    printf("Approximate area: %f\n", global_result); return 0;
}

double Local_trap(double a, double b, int n) {
    double h, x, my_result, local_a, local_b;
    int i, local_n, my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    h = (b - a) / n;
    local_n = n / thread_count;
    local_a = a + my_rank * local_n * h;
    local_b = local_a + local_n * h;
    my_result = (f(local_a) + f(local_b)) / 2.0;
    for (i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    return h * my_result;    //instead of adding it to global_result
}
```

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Area Calculation – v.4

(using *reduction*)

- Instead of using the private `my_result` and the shared `global_result`, the code can use reduction as following:

```
double global_result = 0;    //global_result is shared

# pragma omp parallel num_threads(4) reduction(+:global_result)
global_result += Local_trap(a, b, n); //the + is redundant
    // first, global_result will be private to each thread,
    // on exit, all private results are added(+) into it
```

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Area Calculation – v.4

(using *reduction*)


```
main() {
    double global_result = 0, a=1, b=2;    //global_result is shared
    int n = 12;
    # pragma omp parallel num_threads(4) reduction(+:global_result)
        global_result += Local_trap(...); // or simply =
    printf("Approximate area: %f\n", global_result);
    return 0;
}

double Local_trap(double a, double b, int n) {
    double h, x, my_result, local_a, local_b;
    int i, local_n, my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    h = (b - a) / n;
    local_n = n / thread_count;
    local_a = a + my_rank * local_n * h;
    local_b = local_a + local_n * h;
    my_result = (f(local_a) + f(local_b)) / 2.0;
    for (i = 1; i <= local_n - 1; i++) {
        x = local_a + i * h;
        my_result += f(x);
    }
    return h * my_result; //instead of adding it to global_result
}
```

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Remember!



Global sum


$$8 + 19 + 7 + 15 + 7 + 13 + 12 + 14 = 95$$

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Multiple Cores Forming a Global Sum

- The reduction operator optimises the aggregation of results



Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing

Conclusion/Up Next

- What we covered today (review key concepts):
 - Variable scope (shared, private, firstprivate)
 - Reduction
- Next:
 - Work Sharing

Topic 7: Variable Scope, Reduction

COSC 407: Intro to Parallel Computing