

# COSC 407

## Intro to Parallel Computing

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

1

## Outline (Asynchronous)

### *Previously:*

- Mutual Exclusion (critical, atomic, locks)
- Variable scope (shared, private, firstprivate)
- Reduction, Synchronization (barriers, nowait)

### *Today*

- Work-sharing: parallel for construct
- Data dependency
- Single, master constructs

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

2

# Work-Sharing Constructs

- Within a parallel region, you can decide how work is distributed among the threads.
- ➔ **for** - The for construct is probably the **most important** construct in OpenMP.
- single** - Assigning a task to a single thread
- sections** - Dividing the tasks into sections. Each section is executed by one thread.
- **Implied barrier on exit** (unless **nowait** is used). No implied barrier on entry....
- As they used in parallel region, use existing threads (do not create new threads)

# Parallel For

- Loop iterations are distributed across the thread team (at run time)
  - Loops must have an integer counter variable whose value is inc/dec by a fixed value
  - Reached a specified bound
- Loop variable is explicitly declared to be a private variable (each thread gets own copy)
  - Value at end of loop is undefined (unless set as lastprivate)
  - Recommended that programmers do not rely on OpenMP default rules
    - Explicitly set data sharing attributes
- Output/processing order is non-deterministic

# Parallel For Loop

- Uses a team of threads to execute a **for loop block** following the **for** directive.
  - The system parallelizes the for loop by dividing the iterations of the loop among the threads.

## Syntax (two options)

```
(1) #pragma omp for [clause [clause]...] //does not  
create new threads  
    for(i = start; i <OP> end, incr_expression)
```

The above block must be **placed within a parallel region**.

Use this syntax if your parallel region **includes a for loop and other statements**.

# Parallel For Loop

```
(2) #pragma omp parallel for [clause [clause]...]  
//creates new threads  
    for(i = start; i <OP> end, incr_expression)
```

The above is for creating a parallel block that **ONLY includes** a for loop

**Variable scope:** the loop variable of the for-statement immediately following the for-directive is *private*

- Recall recommend best-practice – declare as private

## WITHOUT the Parallel For Directive

```
#pragma omp parallel num_threads(4)
{
    int i, n = omp_get_thread_num();

    for(i=0; i<4; i++)
        printf("T%d: i=%d\n", n , i);
}
```

- **WITHOUT** the for directive, **EACH** thread runs a copy of the **WHOLE** for loop
  - On 4 threads (T0 to T3) we get 16 print-outs (since each thread executes 4 iterations concurrently with the other threads)
  - Would need to divide up work manually

### Possible Output:

```
T0: i=0
T1: i=0
T0: i=1
T1: i=1
T0: i=2
T1: i=2
T0: i=3
T1: i=3
T2: i=0
T3: i=0
T2: i=1
T3: i=1
T2: i=2
T3: i=2
T2: i=3
T3: i=3
```

## WITH the Parallel For Directive

```
#pragma omp parallel
{
    int i, n;
    #pragma omp for
    for (i = 0; i < 4; i++) {
        n = omp_get_thread_num();
        printf("T%d: i=%d\n", n, i);
    }
}
```

### Possible Output:

```
T0: i=0
T3: i=3
T1: i=1
T2: i=2
```

WITH the for directive, iterations are divided among the threads

- On 4 threads (T0 to T3) we get 4 print-outs (since each thread executes one iteration concurrently with the other threads)
- That is, regardless of the number of threads we will get the exact number of print-outs specified by the for loop.

**Note: order is not preserved!**

# Who Does what?

```
#pragma omp parallel
{
    int i, n;
    #pragma omp for
    for (i = 0; i < 8; i++) {
        n = omp_get_thread_num();
        printf("T%d: i=%d\n", n, i);
    }
}
```


## Possible Output:

```
T1: i=2
T0: i=0
T0: i=1
T2: i=4
T2: i=5
T1: i=3
T3: i=6
T3: i=7
```

Number iterations is divided equally among threads

This is called static scheduling

- T0 : 2 iterations ( i= 0, 1 )
- T1 : 2 iterations ( i= 2, 3 )
- T2 : 2 iterations ( i= 4, 5 )
- T3 : 2 iterations ( i= 6, 7 )



Each thread gets its own loop counter

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

9


# Who Does What?

## cont'd

```
#pragma omp parallel
{
    int i, n;
    #pragma omp for
    for (i = 0; i < 80; i++) {
        //some work
    }
}
```

Number of iterations is divided equally among threads

- T0 : 20 iterations ( i= 0...19 )
- T1 : 20 iterations ( i= 20...39 )
- T2 : 20 iterations ( i= 40...59 )
- T3 : 20 iterations ( i= 60...79 )



Each thread gets its own loop counter

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

10

# Who Does What? *cont'd*

```
//assuming 4 threads
#pragma omp parallel
{
    int i, n;
    #pragma omp for
    for (i = 0; i < 6; i++) {
        n = omp_get_thread_num();
        printf("T%d: i=%d\n", n, i);
    }
}
```

## Possible Output:

```
T0: i=0
T1: i=2
T0: i=1
T1: i=3
T3: i=5
T2: i=4
```

Note: extra iterations are assigned to first few threads

- T0 took 2 iterations ( i= 0, 1 )
- T1 took 2 iterations ( i= 2, 3 )
- T2 took 1 iteration ( i= 4 )
- T3 took 1 iteration ( i= 5 )

0	1	2	3	4	5
T0		T1		T2	T3
i = 0→1		i = 2→3		4→5	5→6

Each thread gets its own loop counter

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

11

# More Practical Examples

- **Image processing:** Converting an RGB image to Grayscale.
  - Assume the luminosity of each grayscale pixel is equal to  $0.21 * \text{Red} + 0.72 * \text{Green} + .07 * \text{Blue}$ . Then:

```
#pragma omp parallel for
for(i=0; i < numPixels; i++)
    gray[i]= (unsigned char)(.21*red[i]+.72*green[i]+.07*blue[i]);
```
- **3D rendering:**
  - Divide image into blocks and iterate over them.
- **Matrix / array operations:**
  - Add two matrixes, find the transpose, etc.
- **Many other applications** that require applying the same action to iteratively in a loop
  - Refer to introduction to parallel computing notes

Topic 8 - Work Sharing (Parallel For, Single)


COSC 407: Intro to Parallel Computing

12

# Area Under the Curve

Revisited

The aim is to compute the area  $S$  (from  $a$  to  $b$ ) under a curve defined by  $f(x)$ .



- One way to do this is to divide  $S$  into a series of trapezoids, find the area of these trapezoids, and then sum.
- The more trapezoids, the better approximation.

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

13

# The Serial Algorithm

Revisited

$$h = \frac{b - a}{n}$$

$$\text{Sum of trapezoid areas} = h \left( \frac{f(a) + f(b)}{2} + f(x_1) + f(x_2) + \dots + f(x_{n-1}) \right)$$

$$x_i = a + i * h$$

```
h = (b - a) / n; //calculate h only once
approx = (f(a) + f(b)) / 2.0; //first and last terms
for (i = 1; i <= n-1; i++){ //remaining terms
    xi = a + i * h;
    approx += f(xi);
}
approx = h * approx; //approximate area
```

Topic 8 - Work Sharing (Parallel For, Single)


COSC 407: Intro to Parallel Computing

14

# Area Calculation

Given the serial code in previous slide, provide the parallel version

- In this version, use the **parallel for** directive
- **Note that when you use parallel-for, it's NOT necessary for n to be evenly divisible by thread\_count**



Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

15

# The Parallel Algorithm

v.5

$$h = \frac{b - a}{n}$$

$$\text{Sum of trapezoid areas} = h \left( \frac{f(a) + f(b)}{2} + f(x_1) + f(x_2) + \dots + f(x_{n-1}) \right)$$

$$x_i = a + i * h$$

```
h = (b - a) / n; //calculate h only once
approx = (f(a) + f(b)) / 2.0; //first and last terms
#pragma omp parallel for reduction(+: approx)
for (i = 1; i <= n-1; i++){ //remaining terms
    xi = a + i * h;
    approx += f(xi);
}
approx = h * approx; //approximate area
```

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

16



# The Complete Program

```
main() {
    double a = 1, b = 2;
    int n = 10, thread_count = 4;
    double global_result = Trap(a, b, n, thread_count);
    printf("Approximate area: %f\n", global_result); return 0;
}

double Trap(double a, double b, int n, int thread_count) {
    double h = (b-a)/n;
    double approx = (f(a) + f(b))/2.0;
    int i;
    #pragma omp parallel for num_threads(thread_count) reduction(+: approx)
    for (i = 1; i <= n-1; i++)
        approx += f(a + i*h);
    return h * approx;
}
```

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

17

# Previous Version without Parallel For

```
main() {
    double global_result = 0, a=1, b=2;    //global_result is shared
    int n = 12;
    # pragma omp parallel num_threads(4) reduction(+:global_result)
        global_result += Local_trap(...);
    printf("Approximate area: %f\n", global_result);
}

double Local_trap(double a, double b, int n) {
    double x, my_approx, my_a, my_b;
    int i, my_n, my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();
    my_n = n / thread_count;
    my_a = a + my_rank * my_n * h;
    my_b = my_a + my_n * h;
    double h = (b - a) / n;
    my_approx = (f(my_a) + f(my_b)) / 2.0;
    for (i = 1; i <= my_n - 1; i++)
        my_approx += f(my_a + i * h);
    return h * my_approx;    //instead of adding it to global_result
}
```

- we need to define range for each thread
- n must be divisible by num\_threads

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

18

# Nested Loops

```
int i, j;
#pragma omp for
for(i = 0; i<5; i++)    //this loop is a parallelized
    for(j=0; j<10; j++) //but this one is not
    ...
```

**Nested for loops:** (In the above example)

- **outer for loop** is parallelized
  - Iterations are divided among threads
- **inner for loop** is not parallelized
  - Each thread will execute all iterations

**Variable scope:**

- The loop counter of the for-statement **immediately following** the for-directive is private for each thread.
  - Hence ( i ) in the above code is private
  - However, ( j ) is **not private** unless it is declared right after the outer for loop, or it is explicitly defined private

# Caution: How and When to Use Parallel For?

Only the following form is legal for parallelized for statement

for	{	index = start ;	index < end	index++	}
		index >= end ;	index <= end	++index	
		index > end	index <= end	index--	
			index >= end	--index	
				index += incr	
				index -= incr	
				index = index + incr	
				index = incr + index	
				index = index - incr	

**Limitations:**

- **index**
  - Must have **integer** or **pointer type** (e.g., it can't be float)
  - Can only be modified by the "increment expression" **in** the **for** statement
- **start**, **end**, and **incr** must:
  - Have a compatible type (e.g., if index is a pointer, then incr must be int)
  - Not change during execution of the loop

## Caution, cont'd

Cannot parallelize:

- **while** or **do-while** loops
- `for (while ; ; ) {...}` //infinite loop
- `for (i=0; i<n; i++){` //cannot determine the  
    `if (...) break` //number of iterations  
}

Watch for **loop-carried dependencies!** (discussed shortly)

- Happen when the calculations in one iteration depends on the data written by other iterations. More details are discussed shortly.
- If they exist, you must do one of two things:
  - Re-write your algorithm to avoid loop-carried dependencies, OR
  - Do not use the parallel-for loops, OR
  - Order your iterations (poor performance)

## Data Dependency

Sequential

```
fibo[ 0 ] = fibo[ 1 ] = 1;  
for (i = 2; i < n; i++)  
    fibo[ i ] = fibo[ i - 1 ] + fibo[ i - 2 ];
```

Parallel

```
fibo[ 0 ] = fibo[ 1 ] = 1;  
#pragma omp parallel for num_threads(8)  
for (i = 2; i < n; i++)  
    fibo[ i ] = fibo[ i - 1 ] + fibo[ i - 2 ];
```

- Output from the parallel code is sometimes correct

1 1 2 3 5 8 13 21 34 55

and sometimes it is wrong

1 1 2 3 5 8 1861489712 -576187344 ...

Q: Explain why!

# Interdependency Among Iterations

- OpenMP **compilers don't check for dependences among iterations** in a loop that's being parallelized with `parallel for`.
- A loop in which the **results of one or more iterations depend on other iterations** cannot, *in general*, be correctly parallelized by OpenMP.

★ **Rule:** *In most cases*, don't use `parallel for` for loops in which the computation of **iterations depends on the data written by other iterations**.

i.e. when one thread reads data written by another thread.

How to address this limitation (again)?

- redesign your algorithm!
- Use 'ordered' clause – has efficiency issues
- Don't use `parallel for`

# Summary of Working With Loops

Basic approach:

- Identify computational intensive loops
- Use one of two techniques:
- **Technique 1:** using `parallel for` or `for` directive
  - Make sure the loops are **iterations-independent** (i.e., don't have loop-carried dependencies)
    - If not, rewrite the algorithm to avoid loop-carried dependencies.
  - Remember that the loop counter of the "for" statement immediately following the `for` directive is private by default.

## Summary of Working With Loops

- **Technique 2:** using `parallel` directive and **manually divide** the range
  - Manually divide the range of iterations and use the thread id (rank) to determine the start, end, and number of iterations for each thread
  - Limitation: number of iterations **must be divisible** by number of threads
- Place appropriate OpenMP directives and test your code
  - Don't forget to protect shared data
  - Use reduction clauses if they simplify the process

## Assigning Work to a Single Thread

### # pragma omp single

Within a parallel region, you can execute a block of code just once by **any** thread in the team.

e.g., when reading in shared input variable or writing single output

*There is an implicit barrier at the end of the "single" region*


```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int main() {
    #pragma omp parallel
    {
        printf("Hi from T%d\n", omp_get_thread_num());
        #pragma omp single
        printf("One Hi from T%d\n", omp_get_thread_num());
    }
    return 0;
}
```

#### Possible Output

```
Hi from T2
Hi from T3
One Hi from T2
Hi from T0
Hi from T1
```

## Assigning Work to a Single Thread, cont'd

```
#pragma omp parallel
{
    int n = omp_get_thread_num();
    printf("T%d:A\n",n);
    #pragma omp single
    printf("T%d:X\n",n);
    printf("T%d:B\n",n);
}
printf("Finished");
```



Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing


27

## Assigning Work to a Single Thread, cont'd

```
#pragma omp parallel
{
    int n = omp_get_thread_num();
    printf("T%d:A\n",n);
    #pragma omp single nowait
    printf("T%d:X\n",n);
    printf("T%d:B\n",n);
}
printf("Finished");
```

### Possible Output

```
T1:A
T0:A
T1:X
T0:B
T1:B
T3:A
T3:B
T2:A
T2:B
Finished
```



Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

28

# Assigning Work to the Master Thread

## #pragma omp master

- Within a parallel region, executes block of code just once by **the master** thread in the team.
- Unlike “single directive”, master directly **does NOT have an implied barrier on exit**.

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main() {
    #pragma omp parallel
    {
        printf("Hi from T%d\n", omp_get_thread_num());
        #pragma omp master
        printf("One Hi from T%d\n", omp_get_thread_num());
    }
    return 0;
}
```

### Possible Output


```
Hi from T2
Hi from T3
Hi from T0
One Hi from T0
Hi from T1
```

# Assigning Work to Master Thread, cont'd

```
#pragma omp parallel
{
    int n = omp_get_thread_num();
    printf("T%d:A\n", n);
    #pragma omp master
    printf("T%d:X\n", n);
    printf("T%d:B\n", n);
}
printf("Finished");
```

### Possible Output

```
T1:A
T1:B
T3:A
T3:B
T0:A
T0:X
T0:B
T2:A
T2:B
Finished
```




(no synchronization **after** 'master' even without nowait)

## Assigning Work to Master Thread, cont'd

```
#pragma omp parallel
{
    int n = omp_get_thread_num();
    printf("T%d:A\n",n);
    #pragma omp master
        printf("T%d:X\n",n);
    #pragma omp barrier
        printf("T%d:B\n",n);
}
printf("Finished");
```

### Possible Output

```
T3:A
T2:A
T0:A
T0:X
T1:A
T3:B
T2:B
-----
T0:B
T1:B
Finished
```



Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

31

## Conclusion/Up Next

- What we covered today (review key concepts):
  - Work-sharing: parallel for construct
  - Data dependency
  - Single, master constructs
- Next:
  - Sections
  - Scheduling Loops (static, dynamic, guided, auto)
  - Ordered Iterations
  - Nested Loops

Topic 8 - Work Sharing (Parallel For, Single)

COSC 407: Intro to Parallel Computing

32