

COSC 407

Intro to Parallel Computing

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

1

Outline (Asynchronous)

Previously:

- Work-sharing: parallel for construct
- Data dependency
- Single, master constructs

Today

- Sections
- Scheduling Loops (static, dynamic, guided, auto)
- Ordered Iterations
- Some examples

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

2



Work-Sharing Constructs

- Within a parallel region, you can decide how work is distributed among the threads.
 - for** - The for construct is probably the most important construct in OpenMP.
 - single** - Assigning a task to a single thread
- ➡ **sections** - Dividing the tasks into sections. Each section is executed by one thread.
 - **Implied barrier on exit** (unless **nowait** is used). No implied barrier on entry....
 - As they used in parallel region, use existing threads (do not create new threads)



Parallel Sections

- Section directive allow us to assign different sections to different threads
- **Each section is executed once.** Each thread executes zero or more sections.
 - A thread may execute zero sections if (# of sections < # of threads)
 - A thread may execute more than one section if it is *fast enough* and the implementation allows it, and/or (# of sections > # of threads)
- **No order!** It is not possible to determine which sections will be executed before which, or if two sections are executed by same thread
 - Therefore, it is important that none of the later sections depends on the results of the earlier ones
- **There is an implicit barrier** at the end of the "sections" region. **No implicit barrier at end of each "section"**
 - There is no implicit barrier on entry



Parallel Sections, Syntax

Within a parallel block, use the following syntax

```
#pragma omp sections [clause [clause]...]  
{  
    #pragma omp section  
        code for section1  
    #pragma omp section  
        code for section2  
    ...  
}
```

If not within a parallel block, use the following syntax:

```
#pragma omp parallel sections [clause [clause]...]  
{  
    #pragma omp section  
        ...  
}
```

The clause could be: **private**, **reduction**, **nowait**


Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

5

Sections – Example 1

```
#pragma omp parallel  
{  
    #pragma omp sections  
    {  
        #pragma omp section  
        {  
            printf("T%d:A \n", omp_get_thread_num());  
            printf("T%d:B \n", omp_get_thread_num());  
        }  
        #pragma omp section  
        {  
            printf("T%d:C \n", omp_get_thread_num());  
            printf("T%d:D \n", omp_get_thread_num());  
        }  
    }  
}
```




Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

6

Sections - Example -1, cont'd

```
#pragma omp parallel sections
{
    #pragma omp section
    {
        printf("T%d:A \n", omp_get_thread_num());
        printf("T%d:B \n", omp_get_thread_num());
    }
    #pragma omp section
    {
        printf("T%d:C \n", omp_get_thread_num());
        printf("T%d:D \n", omp_get_thread_num());
    }
}
```



Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

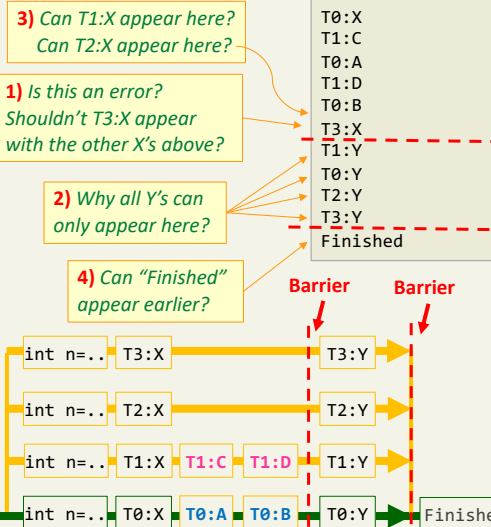
7

Sections - Example - 2

```
#pragma omp parallel num_threads(4)
{
    int n = omp_get_thread_num();

    printf("T%d:X \n", n);
    #pragma omp sections
    {
        #pragma omp section
        {
            printf("T%d:A \n", n);
            printf("T%d:B \n", n);
        }
        #pragma omp section
        {
            printf("T%d:C \n", n);
            printf("T%d:D \n", n);
        }
    }

    printf("T%d:Y \n", n);
}
printf("Finished");
```



Possible Output

```
T1:X
T2:X
T0:X
T1:C
T0:A
T1:D
T0:B
T3:X
T1:Y
T0:Y
T2:Y
T3:Y
Finished
```

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

8

Sections – Example - 2

```

#pragma omp parallel
{
    int n = omp_get_thread_num();
    printf("T%d:X \n", n);

    #pragma omp sections
    {
        #pragma omp section
        {
            printf("T%d:A \n", n);
            printf("T%d:B \n", n);
        }

        #pragma omp section
        {
            printf("T%d:C \n", n);
            printf("T%d:D \n", n);
        }
    }

    printf("T%d:Y \n", n);
}

printf("Finished");

```

Possible Output

```

T1:X
T2:X
T0:X
T1:C
T0:A
T1:D
T0:B
T3:X
T1:Y
T0:Y
T2:Y
T3:Y
Finished

```

COSC 407: Intro to Parallel Computing

9

★ The schedule clause

Determines how iterations are distributed over threads.

```

schedule( static|dynamic|guided|auto
[, chunksize] )

```

Aim: distribute the workload **evenly** so that processors are being used for the same amount of time.

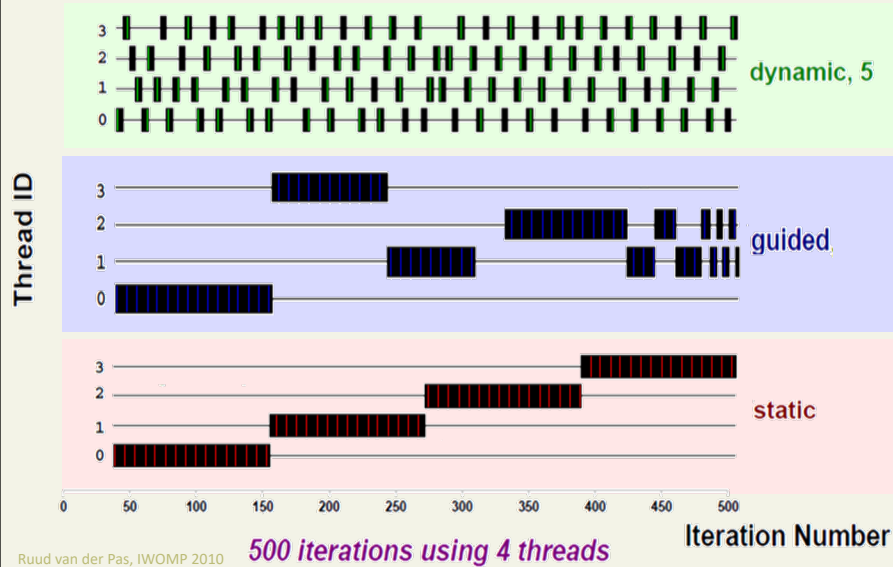
Static is the **default scheduling policy** for most OpenMP implantations

COSC 407: Intro to Parallel Computing

12



Thread Scheduling



Ruud van der Pas, IWOMP 2010

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

13

Static Scheduling

```
schedule( static [,chunksize] )
```

- Before executing the loop, assign the iterations in blocks of size "chunk" over the threads in round-robin fashion.

Default chunk \equiv num_iterations / num_threads

Example: twelve iterations, 0, 1, ..., 11, and 3 threads

```
schedule(static)
```

T0: 0,1,2,3
T1: 4,5,6,7
T2: 8,9,10,11

```
schedule(static,4)
```

T0: 0,1,2,3
T1: 4,5,6,7
T2: 8,9,10,11

```
schedule(static,2)
```

T0: 0,1,6,7
T1: 2,3,8,9
T2: 4,5,10,11

```
schedule(static,1)
```

T0: 0,3,6,9
T1: 1,4,7,10
T2: 2,5,8,11

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

14

Dynamic Scheduling

`schedule(dynamic [,chunksize])`

- Iterations are broken up into chunks of `chunksize`.
- Each thread executes a chunk. Once done, it requests another chunk.
- This keeps going till all iterations are completed.

Default `chunksize` = 1

Example:

```
#pragma omp parallel for schedule(dynamic,2)
for(int i = 0; i<8; i++)
    printf("T%d:%d\n", omp_get_thread_num(),i);
```

Possible outputs

T2:0	T1:4
T1:4	T2:0
T2:1	T0:2
T1:5	T0:3
T0:2	T1:5
T3:6	T0:6
T0:3	T2:1
T3:7	T0:7

Be careful of the **overhead**:

once a chunk is finished, threads need to receive a new iteration counter. To overcome this problem, make sure your chunk size is reasonably large.

Guided/AUTO Scheduling

`schedule(guided | auto [,chunksize])`

Guided

- Similar to dynamic, BUT it **starts with large chunks** then adjusts to smaller chunk sizes if the workload is imbalanced.
 - **Default:**
 - if `chunksize` is unspecified, chunk sizes decrease down to 1.
 - If `chunksize` is specified, it decreases down to `chunksize`
But the last chunk might be smaller than `chunksize`.

Auto

- The compiler and/or the run-time system determine the best schedule. (OpenMP v3+)



When to Use Which?

- **Static:**
 - Use if iterations require roughly the **same amount of time**
 - It requires little overhead (mostly done at compile time)
- **Dynamic**
 - Use if iterations require **different amount of times**.
 - Allows processors the finish first to go after other chunks and hence balance the workload and keep all processors busy
 - It requires more **overhead** than static
- **Guided**
 - Use when the workload increases as we go to higher iterations.
 - i.e. when initial chunks require less time per iteration than later chunks
 - Like dynamic, there is more overhead at runtime

Experiment

We want to parallelize this loop.

```
sum = 0.0;
for(i=0; i<=n; i++)
    sum += f(i);
```

- Assume the time required by **f(i)** is proportional to value of **i**.
- Most OpenMP implementation, static scheduling is used by default.
- A better assignment might be **dynamic** scheduling
 - i.e. cyclic partitioning of the iterations among the threads

Experiment, cont'd

To get a feel for how drastically this can affect performance, assume f is:

```
double f(int i){
    int j, start = i*(i+1)/2, finish = start + i;
    double result = 0.0;

    for(j = start; j<=finish; j++)
        result += sin(j);

    return result;
}
```

- Note that time to run the inner for loop depends on i . for example, $f(2i)$ would take twice the time of $f(i)$

Experiment: Results

$n = 10,000$

one thread (serialized)

run-time = 3.67 seconds.

two threads

Static assignment (default)

run-time = 2.76 seconds

speedup = 1.33

two threads

Dynamic assignment

run-time = 1.84 seconds

speedup = 1.99





The ordered clause

The assignment of iteration chunks to threads is unspecified and hence the **output is usually unordered**.

You can use the **ordered** clause to force **certain events to run in the order of iterations**.

```
#pragma omp for ordered schedule(dynamic)
for(int i=0; i<100; i++) {
    f( a[i] );
    #pragma omp ordered
    g( a[i] );
}
```

f() is done in any order and in parallel


e.g. **f(a[5])** and **f(a[6])** may be done in parallel by two threads

g() is done strictly in order

e.g. assume a thread finishes **f(a[6])** and now wants to proceed to **g(a[6])**. It has to check if **g(a[5])** is finished. If not, the thread **waits** until **g(a[5])** is finished.

The ordered clause

```
#pragma omp for ordered schedule(dynamic)
for(int i=0; i<100; i++) {
    f( a[i] );
    #pragma omp ordered
    g( a[i] );
}
```



- Thin outlines may be done in any order
- Thick outlines must be done in order



Rules and Cost

Rules:

- You can have **exactly one ordered block per an ordered loop**, no less and no more
- The outer for construct must contain the ordered clause

Cost

Ordering comes at an expense of wasting CPU time (waiting for other threads)

Time for (Estimating) PI...

Serial code:

$$\pi = 4 \left[1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

```
int n = 100000000;    //100 millions
double factor = 1, sum = 0;
```

```
for(int k = 0; k < n; k++){
    sum += factor/(2*k+1);
    factor = -factor;
}
```

```
double pi = 4 * sum;
printf("%f", pi);
```

Q: Parallelize the code!

First Piece of PI

First attempt...

Will this code work??

$$\pi = 4 \left[1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

```
int n = 100000000; //100 millions
double factor = 1, sum = 0;
#pragma omp parallel for reduction(+:sum)
for(int k = 0; k < n; k++){
    sum += factor/(2*k+1);
    factor = -factor;
}

double pi = 4 * sum;
printf("%f", pi);
```

Loop-carried
dependency

- a) Yes, it works
- b) No, there is data race
- c) I am super confused!

//possible output: 4.810167

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

25

Second Piece of PI

Second attempt...

How about this solution?

$$\pi = 4 \left[1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

```
int n = 100000000;
double factor, sum = 0;
#pragma omp parallel for reduction(+:sum)
for(int k = 0; k < n; k++){
    if(k%2==0) factor = 1;
    else factor = -1;
    sum += factor/(2*k+1);
}

double pi = 4 * sum;
printf("%f", pi);
```

- a) YES, this is perfect!
- b) No, there is still a problem.
- c) I am still confused!

//possible output: 2.699575

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

26

Third Piece of PI

Third attempt... (Avoiding Data Race)

How about this one now?

$$\pi = 4 \left[1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

```
int n = 100000000;
double factor, sum = 0;
#pragma omp parallel for reduction(+:sum) private(factor)
for(int k = 0; k < n; k++){
    if(k%2==0)    factor = 1;
    else          factor = -1;
    sum += factor/(2*k+1);
}
double pi = 4 * sum;
printf("%f", pi); //output: 3.141593
```

Ensures factor has private scope (WHY?)

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

27

Forth Piece of PI

The Easy Way.... Experiment Contender #1

Contender #1:

Use order clause to force the iterations to run in order

```
int n = 500000000;
double factor = 1, sum = 0;
double time = omp_get_wtime();
#pragma omp parallel for ordered reduction(+:sum)
for(int k = 0; k < n; k++){
    sum += factor/(2*k+1);
    #pragma omp ordered
    factor = -factor;
}
double pi = 4 * sum;
double finish = omp_get_wtime();

printf("pi = %f\ntime = %f ms", pi, 1000*(finish-time));
```

Output
pi = 3.141593
time = 8150 ms

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

28

Third Piece of PI

Experiment Contender #2

Contender #2:

Algorithm redesign (3rd Piece of PI)

```
int n = 500000000;
double factor, sum = 0;
double time = omp_get_wtime();
#pragma omp parallel for reduction(+:sum) private(factor)
for(int k = 0; k < n; k++){
    if(k%2==0) factor = 1;
    else factor = -1;
    sum += factor/(2*k+1);
}
double pi = 4 * sum;
double finish = omp_get_wtime();

printf("pi = %f\ntime = %f ms", pi, 1000*(finish-time));
```

Output
pi = 3.141593
time = 1202 ms

Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

29

Comparison

Contender #1
Using ordered clause

Output
pi = 3.141593
time = 8150 ms

Contender #2
Redesigning the algorithm

Output
pi = 3.141593
time = 1202 ms



Topic 9 - Work Sharing (Sections, Scheduling and Ordered Iterations)

COSC 407: Intro to Parallel Computing

30

Conclusion/Up Next

- What we covered today (review key concepts):
 - Sections
 - Scheduling Loops (static, dynamic, guided, auto)
 - Ordered Iterations
 - Some Examples
- Next:
 - Some More Examples
 - Matrix multiplication
 - Max reduction
 - Asides and Comments on OpenMP