

# COSC 407

## Intro to Parallel Computing

### Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

## Outline

#### *Previously:*

- Error Handling, cudaDeviceSynchronize
- Hardware architecture: sp → SM → GPU
- Thread Organization: threads → blocks → grids
  - Dimension variables (blockDim, gridDim)
- Thread Life Cycle From the HW Perspective
- Kernel Launch Configuration: 1D grids/blocks

#### *Today:*

- Kernel Launch Configuration: nD grids/blocks
- CUDA limits
- Thread Cooperation
- Running Example: Matrix Multiplication

Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu



## Higher Dimensional Grids / Blocks

Remember: choose the breakdown of threads and blocks that **make sense to your problem**.

Example:

- Assume you want to process a 100 pixel x 70 pixel image  
(each 1 thread processes 1 pixel).
- We will have many options, e.g.:

**Option: (1 block/row, 1 thread/pixel)**

A grid of 1x70 blocks (gx=1, gy=70)


each block with 100x1 threads (dx=100, dy=1)



**Another Option (1 block/segment)**

A grid of 10x7 blocks (gx=10, gy=7)

each block with 10x10 threads



Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing



## Higher Dimensional Grids/Blocks

`kernelFunction <<< gridSize , blockSize >>>`

- **gridSize**: dimension and size of the grid in terms of blocks

– could be one of the following:

- `dim3(gx, gy, gz)` → in case of 3D grid  
» Where gx, gy, gz define the three dimensions
- `dim3(gx, gy)` → in case of 2D grid  
» equivalent to `dime3(gx, gy, 1)`
- `dim3(gx)` , or an integer → in case of 1D grid  
» equivalent to `dim3(gx,1,1)` or simply `gx` (the integer)  
» e.g., `dim3(8, 1, 1) = dim3(8) = 8`

- **blockSize**: dimension and size of each block in threads.

- `dim3(bx, by, bz)` → in case of 3D block
- `dim3(bx, by)` → in case of 2D block
- `dim3(bx)` , or an integer → in case of 1D block

Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

# Hello Again....

```
__global__ void hello(){
    printf("Thread(%d,%d,%d) in Block(%d,%d,%d) says:Hello!\n",
        threadIdx.x, threadIdx.y, threadIdx.z,
        blockIdx.x, blockIdx.y, blockIdx.z);
}

int main(){
    hello<<<dim3(2,1,1),dim3(2,1,1)>>>();    // same as hello<<<2,2>>>()
    cudaDeviceSynchronize();                // force the printf() in
                                           // device to flush here

    printf("That's all!\n");
    return 0;
}
```

```
Thread(0,0,0) in Block(0,0,0) says:Hello!
Thread(1,0,0) in Block(0,0,0) says:Hello!
Thread(0,0,0) in Block(1,0,0) says:Hello!
Thread(1,0,0) in Block(1,0,0) says:Hello!
That's all!
```

# Hello Again (same function))

```
__global__ void hello(){
    printf("Thread(%d,%d,%d) in Block(%d,%d,%d) says:Hello!\n",
        threadIdx.x, threadIdx.y, threadIdx.z,
        blockIdx.x, blockIdx.y, blockIdx.z);
}

int main(){
    dim3 gridSize(2,1,1), blockSize(2,1,1);
    hello<<<gridSize, blockSize>>>();
    cudaDeviceSynchronize();    // force the printf() in
                                // device to flush

    printf("That's all!\n");
    return 0;
}
```

```
Thread(0,0,0) in Block(0,0,0) says:Hello!
Thread(1,0,0) in Block(0,0,0) says:Hello!
Thread(0,0,0) in Block(1,0,0) says:Hello!
Thread(1,0,0) in Block(1,0,0) says:Hello!
That's all!
```



## Aside: printf on the kernel?

- Yes, although not a great idea..
  - Specific use cases
- Need to use cudaDeviceSynchronize()
  - Kernel runs asynchronously from host
  - See the code in previous slide

## Computing # of Blocks for 2D Grids

Lets say we have an image of the size WIDTH x HEIGHT

And assume we use 2D blocks of # of threads **TILE\_WIDTH** x **TILE\_HEIGHT**

- E.g. **TILE\_WIDTH** = **TILE\_HEIGHT** = 32, totaling 1024 threads


How do we determine the grid & block organization?

```
//block dimensions
int TILE_WIDTH = 32; //num of threads along x
int TILE_HEIGHT = 32; //num of threads along y
dim3 blockSize(TILE_WIDTH, TILE_HEIGHT);

//grid dimensions
int nblk_x = (WIDTH - 1) / TILE_WIDTH + 1;
int nblk_y = (HEIGHT - 1) / TILE_HEIGHT + 1;
dim3 gridSize(nblk_x, nblk_y);

//launch kernel
kernel<<<gridSize, blockSize>>>(...);
```

```
void kernel(...){
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    ...}
```





# CUDA Limits

For CUDA compute capability 3.0+:

- Within a grid:
  - Total of  $(2^{31} - 1) \times 65535 \times 65535$  blocks
    - Maximum x-dimension of a grid:  $2^{31} - 1$
    - Maximum y- and z- dimension of a grid:  $2^{16} - 1$  (= 65535)
  - That is, **launch as many blocks as you want** (almost)!
- Within a block
  - Maximum total number of threads per block:
    - 1024 (or 512 on older GPUs supporting compute capability < 2)
  - Maximum dimension of a block (# of thread per dimension)
    - x- or y- dimension: 1024 (or 512)
    - z-dimension: 64
- The first assignment on CUDA walks you through this (see last lecture)
- **Check full specs [here](#).**

Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

## FAQs

- Organization of OpenMP threads vs. CUDA threads?
  - OpenMP:
    - number of threads p close to number of processors
  - CUDA:
    - many many threads, organized in 1D, 2D or 3D arrays

Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

Topic 13: CUDA Threads – Part 2

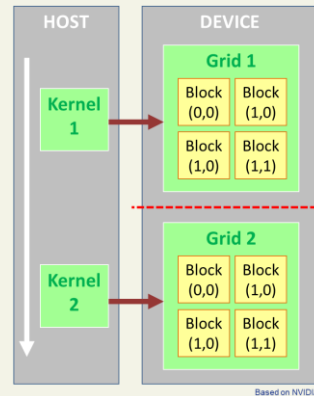
COSC 407: Intro to Parallel Computing



# Threads Cooperation

(more about this later)

- Threads in **same block can cooperate**
  - Synchronize** their execution
  - Communicate via **shared memory**
  - thread/block index is used to assign work and address shared data
- Threads in **different blocks cannot cooperate**
  - Blocks can execute in **any order** relative to other blocks.
  - There is no native way to synchronize all threads in all blocks.
    - To synchronize threads in all blocks**, terminate your kernel at the synchronization point, and then launch a new kernel which would continue with your job

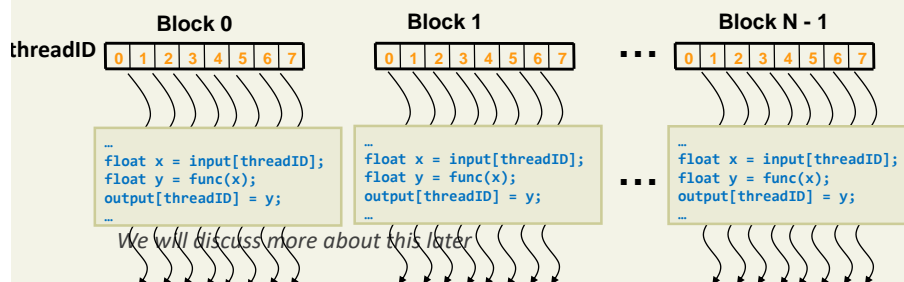


Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

# Threads Cooperation

- For now, all you need to remember is:
  - All threads in **all** blocks run the same kernel.
  - Threads within the **same block** cooperate via shared memory, atomic operations and barrier synchronization.
  - Threads in different blocks **CANNOT** cooperate.



Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

# Matrix Multiplication

- A simple matrix multiplication example that illustrates the basic features of memory and thread management in CUDA programs
  - Assume square matrix for simplicity
  - For now, we will discuss
    - Memory data transfer API between host and device
    - Thread ID usage
  - Later
    - How to speed up performance

Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

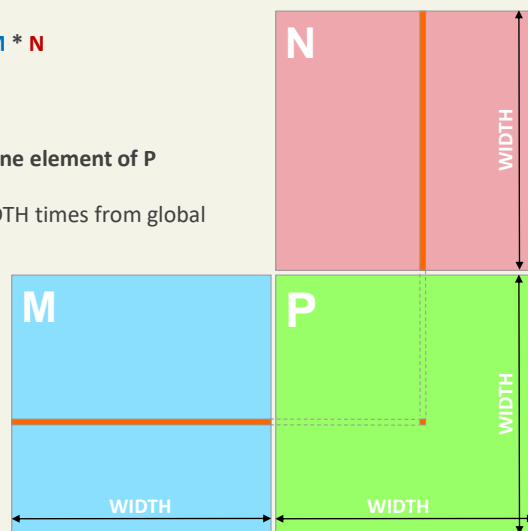
Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

# Programming Model

$$P = M * N$$

- Size is WIDTH x WIDTH
- **Each thread** calculates **one element** of P
- M and N are loaded WIDTH times from global memory




Topic 13: CUDA Threads – Part 2

Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

COSC 407: Intro to Parallel Computing

# How Each Element in P is Computed



Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu


Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

## Serial Code

```
// Matrix multiplication on the (CPU) host in double precision
void MatrixMulSerial(float* M, float* N, float* P, int width)
{
    //for each element Pr,c
    for (int r=0; r<width; r++){ //for each row in P
        for (int c=0; c<width; c++){ //for each column in P
            //Compute the value of Pr,c
            float value = 0;
            for (int k=0; k<width; k++)
                value += M[r*width+k] * N[k*width+c];
            P[r*width+c] = value; //P[r][c]
        }
    }
}
```

```
#pragma omp parallel for private (j, k)
for (r=0; r< n; r++)
    for (c=0; c<n; c++)
        for (k=0; k<m; k++)
            P[r][c] += M[r][k] * N[k][c];
```



Topic 13: CUDA Threads – Part 2

Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

COSC 407: Intro to Parallel Computing



# Processing 2D arrays in C

Review

Let's say we have two 2D arrays that we want to process them in loops (e.g. initialize to 0's)

**Serial Code:**

```
int A[10][10];
int* B = malloc(100 * sizeof(int));

for (int r = 0; r < 10; r++){
    for (int c = 0; c < 10; c++) {
        A[r][c] = 0; // this is ok
        B[r][c] = 0; // ERROR!
    }
}
```

You cannot use [r][c] with dynamically allocated arrays (B can only be a pointer or a vector).

**Solution:** use row-major format!

Topic 13: CUDA Threads – Part 2


COSC 407: Intro to Parallel Computing

# Memory Layout: Matrix in C

Review

M <sub>0,0</sub>	M <sub>1,0</sub>	M <sub>2,0</sub>	M <sub>3,0</sub>
M <sub>0,1</sub>	M <sub>1,1</sub>	M <sub>2,1</sub>	M <sub>3,1</sub>
M <sub>0,2</sub>	M <sub>1,2</sub>	M <sub>2,2</sub>	M <sub>3,2</sub>
M <sub>0,3</sub>	M <sub>1,3</sub>	M <sub>2,3</sub>	M <sub>3,3</sub>

2D arrays in C are stored like this in the memory

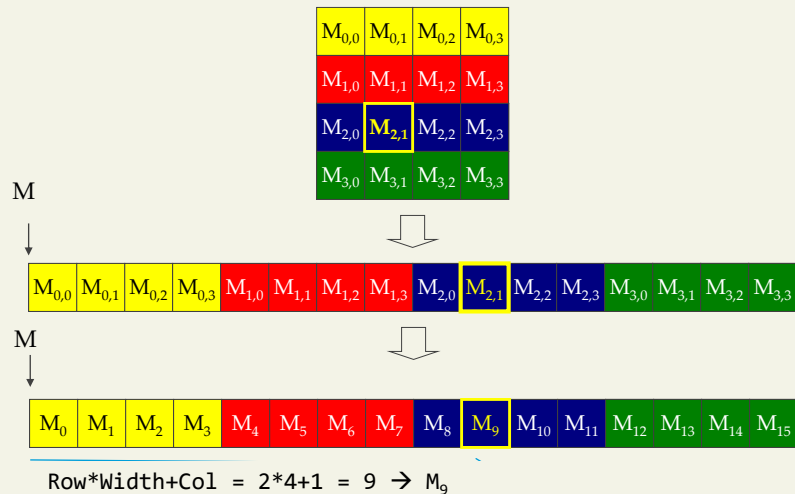


Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

# Row-Major Layout in C/C++



$$\text{Row} * \text{Width} + \text{Col} = 2 * 4 + 1 = 9 \rightarrow M_9$$

$$M[r][c] = M[r * \text{Width} + c]$$

# Processing 2D arrays in C

Review cont.

Let's say we have two 2D arrays that we want to process them in loops (e.g. initialize to 0's)

## Serial Code:

```
int A[10][10];
int* B = malloc(100 * sizeof(int));

for (int r = 0; r < 10; r++){
    for (int c = 0; c < 10; c++) {
        A[r][c] = 0;           // this is ok
        B[r*10+c] = 0;         // this is ok now!
    }
}
```


# Parallel Code: Using *One* Block

## Basic Idea

- Only **ONE** block used to compute the output matrix P
- Each thread computes one element of P as follows:
  - Loads a row of matrix M
  - Loads a column of matrix N
  - Perform one multiply and addition for each pair of M and N elements
  - Compute and stores the result on an off-chip memory (DRAM)

## Limitation:

- Size of P is limited to 32x32
  - i.e. the number of threads allowed in a thread block.



Slide materials based on: 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

Topic 13: CUDA Threads – Part 2


COSC 407: Intro to Parallel Computing

# Parallel: Kernel – *One* Block

// Matrix multiplication kernel – each thread computes one P element

```
__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int width){
    //find index of Pr,c element
    int r = threadIdx.y + blockIdx.y * blockDim.y;
    int c = threadIdx.x + blockIdx.x * blockDim.x;
    //compute P's element
    if(r<width && c<width){
        float value = 0;
        for (int k=0; k<width; k++)
            value += d_M[r*width+k] * d_N[k*width+c];
        d_P[r*width+c] = value;
    }
}
```

Also ok to use  
`int r = threadIdx.y;`  
`int c = threadIdx.x;`  
 But it is better to use the general formula ( here, we use only one block, and thus blockIdx = 0). **WHY BETTER?**



## Parallel : Host – *One* Block

```
void MatrixMulOnDevice(float* M, float* N, float* P, int width)
{
    int size = width * width * sizeof(float);
    float *d_M, *d_N, *d_P;
    //1) Allocate M, N, P on device memory. Copy M,N to device
    cudaMalloc(&d_M, size);
    cudaMalloc(&d_N, size);
    cudaMalloc(&d_P, size);
    cudaMemcpy(d_M, M, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_N, N, size, cudaMemcpyHostToDevice);
    //2) Kernel invocation code
    dim3 blockSize(width, width);
    MatrixMulKernel<<<1, blockSize>>>(d_M, d_N, d_P, width);
    //3) Read P from the device
    cudaMemcpy(P, d_P, size, cudaMemcpyDeviceToHost);
    //4) Free device matrices
    cudaFree(d_M); cudaFree(d_N); cudaFree(d_P);
}
```


Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

## Using *Multiple* Blocks

- We saw that using only **one block** has a **serious limitation**: size of matrix limited by 1024.
- Also, you are not fully using your GPU
- **Solution**: use multiple blocks
  - We shall apply the method explained previously
- More on this next day



Slide materials based on, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign  
© David Kirk/NVIDIA and Wen-mei W. Hwu

Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

# Remember...

*Why we need to divide threads into blocks with the grid?*

- To make thread organization better fit the problem
  - e.g., 2D blocks for 2D images.
- To satisfy CUDA limits (only 1024 threads per block)
  - We also need to avoid GPU hardware limits
    - **For example, G80** has 16 SMs.
      - Each **SM** can process up to **8 blocks** at a time and up to **768 threads** at a time (*more later*)
- To exploit the GPU full power
  - E.g., one block means one SM is functioning and remaining are not
- To allow for threads communication at different levels
  - Threads within same block have “shared memory” and can sync. Threads in different blocks cannot sync (at least directly) and can only share data through the global memory.
  - **More about this next...**

Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing

# Summary

## **Today:**

- Kernel Launch Configuration: nD grids/blocks
- CUDA limits
- Thread Cooperation
- Running Example: Matrix Multiplication

## **Next:**

- Tiling
- CUDA Scalability
- Thread Scheduling on the H/W: Thread Lifecycle
  - zero-overhead and latency tolerance
- GPU limits
- CUDA Memories Types (and Performance)
- Example: Improving Performance of Matrix Multiplication

Topic 13: CUDA Threads – Part 2

COSC 407: Intro to Parallel Computing