

COSC 407

Intro to Parallel Computing

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

1

Outline

Previous pre-recorded lecture (Students' led Q/As):

- CUDA basics: program structure
- Useful Built-in CUDA functions
- Function Declarations (global, device, host)

Today:

- Error Handling, cudaDeviceSynchronize
- Hardware architecture: sp → SM → GPU
- Thread Organization: threads → blocks → grids
 - Dimension variables (blockDim, gridDim)
- Thread Life Cycle From the HW Perspective
- Kernel Launch Configuration: 1D grids/blocks

Next Lecture:

- Kernel Launch Configuration: nD grids/blocks
- CUDA limits
- Thread Cooperation
- Running Example: Matrix Multiplication

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

2



Memory

CPU and GPU have separate memory spaces

- Need to move data to device (GPU) if it is processed there
- Need to move results back to CPU memory
- Functions: `cudaMalloc`, `cudaFree`, `cudaMemcpy`

Pointers

- Hold memory addresses in either CPU or GPU memory
- Can't differentiate CPU pointers from GPU pointers by just checking their values.
 - There, you must use pointers in their appropriate locations.
 - Dereferencing CPU pointer in kernel will likely crash
 - Dereferencing GPU pointer host code will likely crash

Error Handling

CUDA has two sources of errors :

(1) Errors from CUDA API

E.g. cannot allocate memory space on the device

(2) Errors from CUDA Kernel

i.e. errors that happen inside your kernel code.



Handling CUDA API Errors

- CUDA API functions return an error code of type `cudaError_t`
- For example:
 - `cudaSuccess` (=0, if no problems)
 - `cudaErrorMemoryAllocation` (=2, if cannot allocate memory)
 - Other error codes (positive values) are possible
 - see [here](#) for the full list.

Such errors should be handled using some extra code. For example:

```
cudaError_t err = cudaMalloc(&d_a, num_bytes);
if (err != cudaSuccess) {
    printf("Can't allocate CUDA Memory");
    ...//more code to handle error
} else {...}
```

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

5



Handling CUDA API Errors

A better IDEA: to avoid repeatedly writing if statements after each CUDA call, you can define and use a macro as following:

```
#define CHK(call) { \
    cudaError_t err = call; \
    if (err != cudaSuccess) { \
        printf("Error%d: %s:%d\n", err, __FILE__, __LINE__); \
        printf(cudaGetErrorString(err)); \
        cudaDeviceReset(); \
        exit(1); \
    } \
}
```

destroys and clean up all resources associated with the current device in the current process immediately

Then use this macro whenever you call a CUDA API function

```
CHK( cudaMalloc(&d_a, num_bytes) );
```

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

6



Handling CUDA Kernel Errors

- The other type of errors is the one that happens during the execution of YOUR kernel function
- You can check for this error as follows

```
Kernel<<<...>>>();           //call kernel
CHK(cudaGetLastError());      //1
CHK(cudaDeviceSynchronize()); //2
```

- Statement #1 will **check for kernel launch errors**
 - e.g. too many threads per block
 - CUDA runtime maintains an error variable that is overwritten each time an error occurs. **cudaGetLastError()** returns the value of this variable and resets the variable to **cudaSuccess**.
- Statement #2 will block the host until GPU is done
 - Any asynchronous error is returned by (cudaDeviceSynchronize)

More details: <https://devblogs.nvidia.com/how-query-device-properties-and-handle-errors-cuda-cc/>

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

7



cudaDeviceSynchronize()

- CUDA functions **and** host code are **asynchronous**
 - i.e. they return control to the calling CPU thread before they finish their work
- **cudaDeviceSynchronize()** can be used to block the calling CPU thread until all CUDA calls made by this thread are finished
- Example use: time your kernel
 - (must include time.h and cuda lib)

```
//on host
double t = clock();
Kernel<<<...>>>();
cudaDeviceSynchronize()
t = (clock()-t)/CLOCKS_PER_SEC;
```

Synchronization is expensive, so don't overuse it!

Topic 13: CUDA Threads

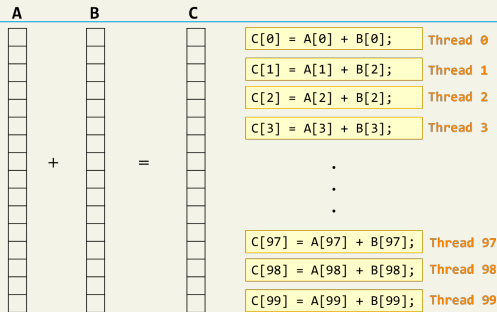
COSC 407: Intro to Parallel Computing

8

Adding Vectors: Revisited

- You saw before that we usually assign a thread to process each element in an array.
- Assigning threads to *vector* elements was easy

```
__global__ void vec_add(float *A, float *B, float* C, int N) {  
    int i = threadIdx.x;  
    if (i < N) C[i] = A[i] + B[i];  
}
```



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

9

Typical GPU Program

Host code (running on CPU)

1. Allocate space on GPU
2. Copy CPU data to GPU
3. Launch kernel function(s) on GPU

define launch-configuration before that.

4. Copy results from GPU to CPU
5. Free GPU memory

Kernel code (running on GPU)

- Write kernel function as if it will run on a single thread
 - Use IDs to identify which **piece of data** is processed by this thread
 - Remember that this SAME kernel function is executed by many threads
- Parallelism of threads is expressed in the host code

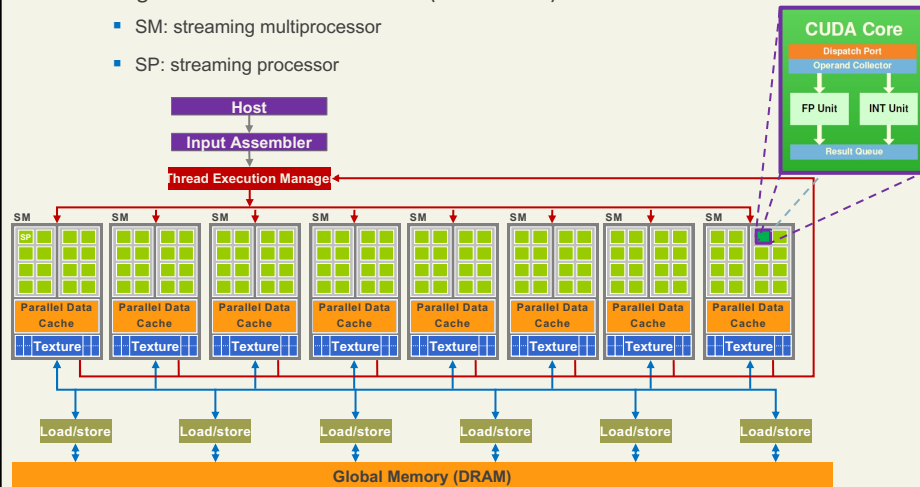
Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

10

GPU Design

- Massively threaded, sustains 1000s of threads per app
- The figure: 8 SMs x 16 SP = 128 SPs (CUDA cores)
 - SM: streaming multiprocessor
 - SP: streaming processor



Topic 13: CUDA Threads

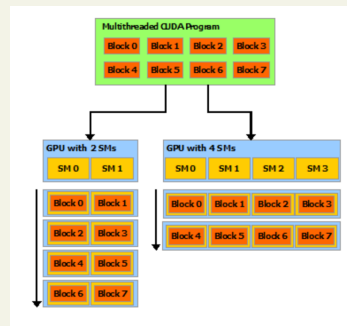
COSC 407: Intro to Parallel Computing

11



GPU Design

- A scalable array of multithreaded *Streaming Multiprocessors* (SMs)
- A multithreaded program is partitioned into blocks of threads that execute independently from each other
 - GPU with more multiprocessors will automatically execute the program in less time than a GPU with fewer multiprocessors.



<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

12



GPU Design

- When host invokes a kernel grid
 - Blocks of the grid are enumerated
 - Distributed to multiprocessors with available execution capacity
- The threads of a thread block execute concurrently on one multiprocessor
 - Multiple thread blocks can execute concurrently on one multiprocessor
- As thread blocks terminate
 - New blocks are launched on the vacated multiprocessors
- Designed to execute hundreds of threads concurrently. To manage such a large amount of threads (*SIMT*)

13




Threads Organization: Basics

On SOFTWARE (code) side:

- **Threads** are grouped into **Blocks**
 - All threads in a block execute the same kernel program (*SPMD*)
- **Blocks** are grouped into a **Grid**
- **IDs:**
 - Each **thread** has a unique ID *within a block*
 - Each **block** has a unique ID *within a grid*

On HARDWARE side:

- Each **block** runs on one **SM**.
 - An **SM** might run **more than one block**
- Each thread runs on an **SP** (within an **SM**)
 - An **SP** can only run one thread at any time
 - Might run many successive threads.
- *More about this later*



14

Device Properties

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"
#include <stdio.h>

int main()
{
    //just to check
    cudaDeviceProp prop;
    int count;
    cudaGetDeviceCount(&count);
    for (int i = 0; i < count; i++)
    {
        cudaGetDeviceProperties(&prop, i);
        ...
        //examine members of the struct
    }
}
```

```
----- General Information for device 0 -----
Name: Tesla K80
Compute capability: 3.7
Clock rate: 823500
Device copy overlap: Enabled
Kernel execution timeout: Disabled
----- Memory Information for device 0 -----
Total global mem: 11996954624
Total constant Mem: 65536
Max mem pitch: 2147483647
Texture Alignment: 512
----- MP Information for device 0 -----
Multiprocessor count: 13
Shared mem per mp: 49152
Registers per mp: 65536
Threads in warp: 32
Max threads per block: 1024
Max thread dimensions: (1024, 1024, 64)
Max grid dimensions: (2147483647, 65535, 65535)
```

Topic 13: CUDA Threads



COSC 407: Intro to Parallel Computing

15



Threads in a Block

- Threads in a block are organized in **1D, 2D, or 3D** array of threads.
- Built-in **ID variables**
 - `threadIdx.x`
 - `threadIdx.y`
 - `threadIdx.z`
- Thread IDs are **unique within a block**



```
threadIdx.x = 2
threadIdx.y = 1
threadIdx.z = 1
```


Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing


16

Why 1/2/3-D Organization?

Simplifies memory addressing when processing **multidimensional data**




1D threads are most suitable for processing **vectors**



2D threads are most suitable for 2D **arrays** (e.g. images)

3D threads are most suitable for 3D **arrays** (e.g. 3D environments)



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

17

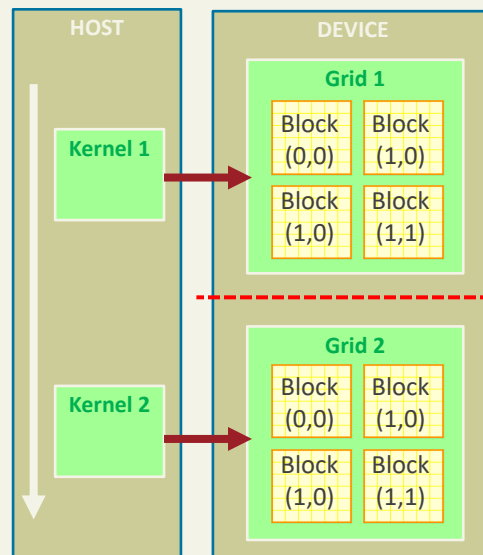


Blocks in a Grid

- Kernel code (on host) may initiate one or more blocks, each with many threads.

```
__global__ kernel1(..) {..}
...
kernel1<<<gridSize,blockSize>>> (...);
...
```

- All blocks for a given kernel belong to a grid
- All blocks in a grid **must finish before the next kernel runs**.
 - Synchronization point!
- Remember that each block runs on one SM



Based on NVIDIA

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

18



Blocks in a Grid


- Blocks in a grid are organized in 1D, 2D, or 3D array of blocks.

- Built-in ID variables

`blockIdx.x`

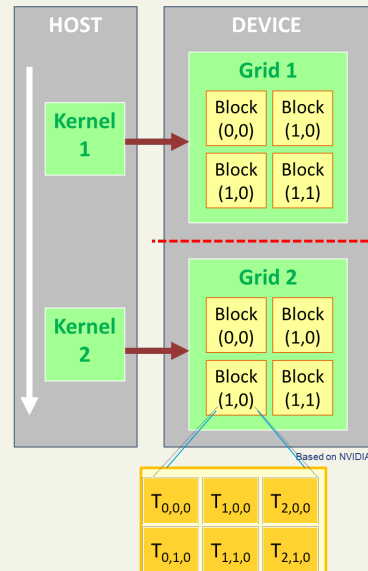
`blockIdx.y`

`blockIdx.z`



`blockIdx.x = 1`
`blockIdx.y = 1`
`blockIdx.z = 1`

- Block IDs are **unique within a grid**



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

19



Dimension Variables

A dimension variable holds the number of elements over this dimension

- Dimensions may be unique for each grid and are set at launch time
 - cannot change a kernel's dimensions once it is launched.

- Built-in **dimension** variables

`blockDim.x`, `blockDim.y`

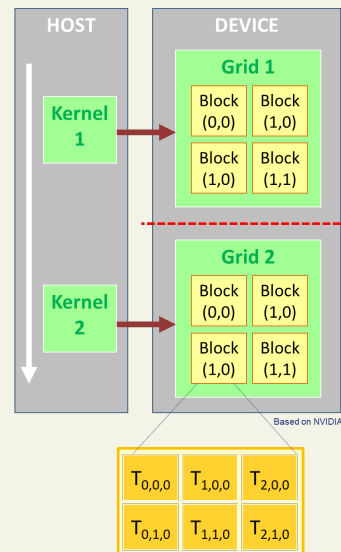
`blockDim.z`

`gridDim.x`, `gridDim.y`,

`gridDim.z`

- E.g. Grid 2 in the figure:

`gridDim.x = 2` `blockDim.x = 3`



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

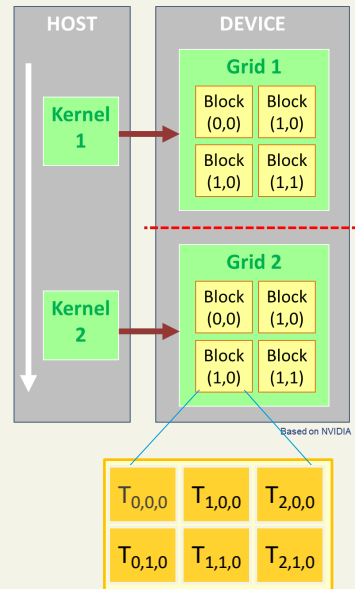
20

Thread Life Cycle in HW

Incomplete version

1. Grid is launched
`kernelFoo<<<gridSize,blockSize>>> (...);`
2. Blocks are distributed to SM
— Potentially more than one Block per SM
3. Each SM launches the threads in its block.
— One thread per core (SP)
4. As Blocks complete, resources are freed.

Note: this is not the complete lifecycle. We are still missing the “warps” which are discussed later.



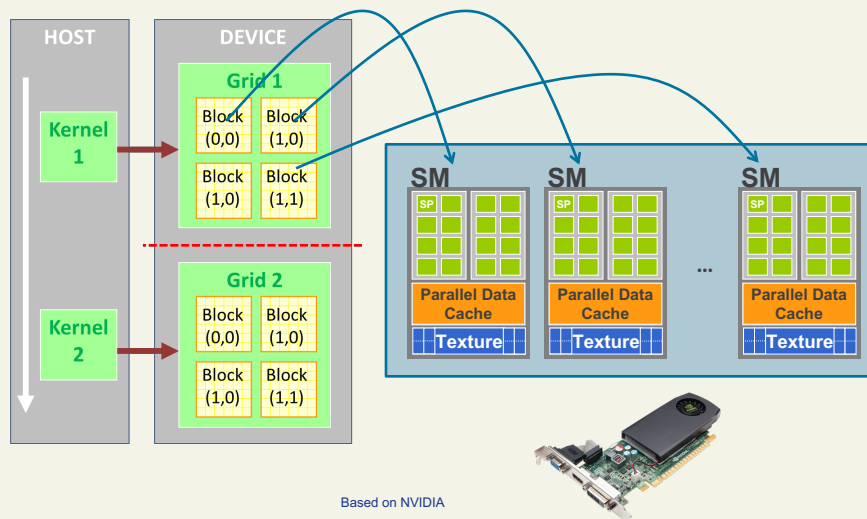
Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

21

Thread Life Cycle in HW

Incomplete version



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

22



Kernel Launch Configuration

- From the Vector Addition Example


```
vectorAdd<<<1, N>>>(...);
```

 - Statement tells the GPU to launch **N** threads on **1** block

- The general format:

```
kernelFunc<<<gridSize, BlockSize>>>
```


- You can:
 - Run **as many blocks** at once (all belong to the same grid)
 - Each block can have a **maximum of**:
 - 1024** threads on newer GPUs.
 - 512** threads on older GPUs
- We **cannot** specify which block runs before which.

Kernel Launch Configuration

- You should choose the breakdown of threads and blocks that **make sense to your problem**.

Example: For a vector, choose 1D setup with options

`KernelFunc<<< 1, 30 >>>(...);`



30 threads: 1 Block with 30 threads

`KernelFunc<<< 3, 10>>>(...);`



30 threads: 3 Blocks, each with 10 threads

- Dimensionality of above example:** 1D blocks and 1D grid
- x-dimension** is used by default for 1D items
 - Can define **higher dimensionality** using `dim3`. (more about this later)

Remember, each block is assigned to one SM. If you want to fully use the GPU, then **#blocks should be \geq # of SMs**

Vector Addition: Revisited

```
__global__ void vectorAdd(int* a, int* b, int* c, int
n) {
    int i = threadIdx.x;
    if(i<n)
        c[i] = a[i] + b[i];
}
```

Only x-dim is used

```
int main() {
    int *a, *b, *c, *d_A, *d_B, *d_C;
    //...allocate space on CPU and GPU
    //...initialize a,b
    //...copy a,b to GPU at d_A, d_B
    //launch the kernel
    vectorAdd <<<1,N>>> (d_A, d_B, d_C, N);
    //...results back from d_C to c
    //...free up memory
    return 0;
}
```

This means 1 grid with 1 block running on 1 SM, and N threads organized in 1D array (over the x-dimension only)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2010, ECE 408, University of Illinois, Urbana-Champaign

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

25

Computing Array Index for Each Thread



foo<<< 1, 30 >>>(...);



1 Block with 30 threads

```
void foo(...){
    int i = blockIdx.x * blockDim.x
    ... + threadIdx.x
}
```

Example: computing i for $v[2]$, $v[5]$



$i=0*4+2=2;$

$i=0*4+5=5;$

Topic 13: CUDA Threads



foo<<< 3, 10>>>(...);



3 Blocks, each with 10 threads

```
void foo(...){
    int i = blockIdx.x * blockDim.x
    ... + threadIdx.x
}
```

Example: computing i for $v[2]$, $v[5]$



$i=0*4+2=2;$

$i=1*4+1=5;$

COSC 407: Intro to Parallel Computing

29

Computing Array Index for Each Thread

- The general formula to compute the thread index:

```
int x = blockIdx.x * blockDim.x + threadIdx.x;
int y = blockIdx.y * blockDim.y + threadIdx.y;
int z = blockIdx.z * blockDim.z + threadIdx.z;
```


Use above formulas *in the kernel function* to identify which data element is accessed by each thread.

- Example:** using the formulas above, compute the (x,y) of the highlighted element; i.e. confirm that thread $T_{1,1,0}$ in block $B_{1,0,0}$ will be only accessing `img[1][4]`

Answer: For this thread:

$x = 1*3+1 = 4$ //col

$y = 0*2+1 = 1$ //row



Topic 13: CUDA Threads



COSC 407: Intro to Parallel Computing

30

Processing 100x70 Picture

```
__global__ void PicKrn1(float* d_Pin, float* d_Pout, int w, int h){
    // Calculate row # of the d_Pin and d_Pout element to process
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    // Calculate column # of the d_Pin and d_Pout element to process
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    // each thread computes one element of d_Pout if in range
    if((y<h)&&(x<w)) d_Pout[y * w + x] = f(d_Pin[y * w + x]);
}
```

Using 16x16 blocks



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing


31

Processing 100x70 Picture

```
__global__ void PicKrn1(float* d_Pin, float* d_Pout, int w, int h){
    // Calculate row # of the d_Pin and d_Pout element to process
    int y = blockIdx.y;
    // Calculate col # of the d_Pin and d_Pout element to process
    int x = threadIdx.x;
    // each thread computes one element of d_Pout if in range
    if((y < h) && (x < w)) d_Pout[y * w + x] = f(d_Pin[y * w + x]);
}
```

Using 100x1x1 blocks

Only use red code if you are sure you have a 1D block.



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

32


Processing 100x70 Picture

```
__global__ void PicKrn1(float* d_Pin, float* d_Pout, int w, int h){
    // Calculate row # of the d_Pin and d_Pout element to process
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    // Calculate col # of the d_Pin and d_Pout element to process
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    // each thread computes one element of d_Pout if in range
    if((y < h) && (x < w)) d_Pout[y * w + x] = f(d_Pin[y * w + x]);
}
```

Using 100x1x1 blocks

Using the General Formula gives same output (and it's preferred)

Try it: since is a 1-row block, put *threadIdx.y = 0*, *blockDim.y = 1*, *blockIdx.x = 0* and the code will be the same as the previous slide, **except that** it also works if you decide to change dimensionality



Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

33

Choosing Launch Config

(A) 1D Grids / Blocks

- Assume you know total number of **threads (N)** needed
 - Should be equal to the number of data elements
- How do you determine the launch configuration?

`KernelFunc<<< ???, ??? >>>`

Steps:

1. Choose the number **threads per block (nthreads)**.
2. Compute the **number of blocks** as follows:

$$\text{nblocks} = (\text{N}-1)/\text{nthreads} + 1$$

Note (again): if you want to fully use the GPU, then

- #threads per block should be large (\geq #SPs per SM)
- #blocks should be \geq # of SMs

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

34

Launch Configuration Examples

Assume we choose `nthreads = 256` (i.e. #threads per block)

- # of array elements **N** = 200
 - `nblocks = 199/256 + 1 = 1` → total # threads = 256
- # of array elements **N** = 256
 - `nblocks = 255 / 256 + 1 = 1` → total # threads = 256
- # of array elements **N** = 400
 - `nblocks = 399 / 256 + 1 = 2` → total # threads = 512

Note: use `if(i < n)` to discard the extra threads

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

35


Example: Vector Addition

```
__global__ void vectorAdd(int* a, int* b, int* c, int n) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if(i < n) c[i] = a[i] + b[i];
}
```

Now you know why

```
int main() {
    int *a, *b, *c, *d_A, *d_B, *d_C;
    //...allocate space on CPU and GPU
    //...initialize a,b
    //...copy a,b to GPU at d_A, d_B
    //launch the kernel - Assume array has N elements
    int nblocks = (N-1)/512+1; //512 threads per block
    vectorAdd <<<nblock,N>>> (d_A, d_B, d_C, N);
    ...
}
```

Grid



1D array of blocks, each having 1D array of threads

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

36

Vector Addition: Full Code

Rewrite the serial program below so that vectorAdd runs on the GPU with 4 blocks each having 256 threads



SERIAL CODE

```
#define N 1024
void vectorAdd(int* a, int* b, int* c, int n) {
    int i;
    for (i = 0; i < n; i++)
        c[i] = a[i] + b[i];
}

int main() {
    int *a, *b, *c, i;
    a = (int*) malloc(N * sizeof(int)); // create three arrays
    b = (int*) malloc(N * sizeof(int));
    c = (int*) malloc(N * sizeof(int));

    for(i=0;i<N;i++) a[i] = b[i] = i; // initialize a and b for testing
    vectorAdd(a, b, c, N); // serial vector addition
    for(i=0;i<10;i++) printf("c[%d] = %d\n", i, c[i]);
    free(a); free(b); free(c); // free up memory taken by a,b,c
    return 0;
}
```


Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

40

Vector Addition : Parallel Code

Step1: parallelizing the function



```
__global__ void vectorAdd(int* a, int* b, int* c, int n) {
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if(i < n)
        c[i] = a[i] + b[i];
}
```

Vectors are 1D,
so we assume
we launch a 1D
grid of threads

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

41

Vector Addition: Parallel Code

Step2: modify the main method

```
int main() {
    int *a, *b, *c, i; //pointers to host memory + loop counter
    int *d_A, *d_B, *d_C; //pointers to device memory

    a = (int*) malloc(N * sizeof(int)); //allocate space on host
    b = (int*) malloc(N * sizeof(int));
    c = (int*) malloc(N * sizeof(int));

    //allocate space on device
    cudaMalloc(&d_A, N * sizeof(int));
    cudaMalloc(&d_B, N * sizeof(int));
    cudaMalloc(&d_C, N * sizeof(int));

    for(i=0; i<N; i++) a[i]=b[i]=i; //initialize a,b (for testing)
    //copy data (i.e. a and b) from host to device
    cudaMemcpy(d_A, a, N * sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, b, N * sizeof(int), cudaMemcpyHostToDevice);
}
```

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

42

Vector Addition: Parallel Code

```
//specify threads configuration & pass pointers to device memory
int nThreads = 256;
int nBlocks = N/nThreads
if (N%256) nBlocks++;

vectorAdd<<<nBlocks, nThreads>>>>(d_A, d_B, d_C, N);

//copy results from device to host
cudaMemcpy(c, d_C, N * sizeof(int), cudaMemcpyDeviceToHost);

for(i = 0; i < 10; i++) // print first 10 elements(for testing)
    printf("c[%d] = %d\n", i, c[i]);

free(a);free(b);free(c); // free the host memory taken by a,b,c

//free device memory
cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);

return 0;
}
```

Another way to specify the # of blocks given # of threads and # of threads/block

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

43

Summary: How Many Blocks?

Given that **nthreads** is # threads per block

Method1:

```
nblocks = (N-1)/nthreads + 1
dim3 gridSize(nblocks, 1, 1);
dim3 blockSize(nthreads, 1, 1);
```

Method2:

```
nblocks = N/nthreads;
if(N%256) nblocks++; //if there is remainder, add one more block
dim3 gridSize(nblocks, 1, 1);
dim3 blockSize(nthreads, 1, 1);
```

Topic 13: CUDA Threads

COSC 407: Intro to Parallel Computing

44

Summary

Previous pre-recorded lecture (Students' led Q/As):

- CUDA basics: program structure
- Useful Built-in CUDA functions
- Function Declarations (global, device, host)

Today:

- Error Handling, cudaDeviceSynchronize
- Hardware architecture: sp → SM → GPU
- Thread Organization: threads → blocks → grids
 - Dimension variables (blockDim, gridDim)
- Thread Life Cycle From the HW Perspective
- Kernel Launch Configuration: 1D grids/blocks

Next Lecture:

- Kernel Launch Configuration: nD grids/blocks
- CUDA limits
- Thread Cooperation
- Running Example: Matrix Multiplication